

The Autonomous Boardroom

Responsible AI Governance for Multi-Agent Strategic Cascades in FMCG

Case Company: SA Beverages (Anonymized) | Growth Vector: Diversification
New product line: 'Savories' baked snacks | Core business: packaged beverages, bottling and distribution

Shivam Ahuja

Contents

- Executive Summary
- FMCG Problem Statement & Business Case (Phase 1)
- Top-Management Debate, Design, Execution & Mandate (Phase 2)
- Middle-Management Debate, Design, Execution & Defense (Phase 3)
- NLP Methodology & Cross-Level Vulnerability Diagnostic (Phase 4)
- Responsible AI Governance Framework & Recommendations (Phase 5)
- References

Executive Summary

SA Beverages, a leading Indian packaged-beverages company, faces decelerating growth, from roughly 25 percent in 2023 to 2024 down to about 8 percent in 2024 to 2025, alongside heavy dependence on a single franchise partner. Its CEO wants to diversify into packaged savoury snacks under a new sub-brand, Savories. Across FMCG, decisions exactly like this are beginning to be piloted with autonomous multi-agent AI. This paper stress-tests that idea. It runs the decision through two AI debates, measures how they behave, and turns the differences into a governance framework the leadership can adopt. All figures are drawn from a real listed beverage major and the company identity is fictionalized.

Two debates were run on the same model, Claude Opus 5. The first was a fully autonomous board of five executives that argued freely and adversarially, and it endorsed the plan only conditionally, at 6.1 out of 10. The second was a human-led cascade in which the author played the Product Head defending the board's mandate to four functional leads, under one deliberate constraint, the team could raise execution concerns but could not reopen the strategy. That suppressed-pushback team rated the same plan far more favourably, at 7.8 out of 10.

An NLP diagnostic on the two transcripts showed, across five independent measures, that the suppressed-pushback debate behaved differently from the adversarial one in consistent and worrying ways. The team was far more positive and agreeable than the board, it echoed its leader roughly three times as strongly, its four voices converged into one about twice as much, its strategy vocabulary was halved as the conversation narrowed to execution, and yet it rated its own plan higher. In short, the debate with less scrutiny showed more confidence.

The central lesson is that a smooth, agreeable AI debate is not a safe one, it is an unmonitored one. Contestation is the control. The paper therefore recommends the CANDOR framework, five controls that each neutralise one measured vulnerability, from a mandated challenger role to authority-blind reasoning, a preserved minority report, a protected strategy checkpoint, and confidence that is discounted by the scrutiny it received. A sixth element turns the diagnostic metrics into a live governance dashboard, calibrated so that the adversarial board is the healthy band and the suppressed team is the danger band. Human attention is then tiered to the stakes of each decision rather than spread thinly.

For SA Beverages the framework is cheap to start, because the monitoring is simply the same analysis run automatically on any multi-agent transcript, and it can be piloted on the Savories decision itself. More broadly, the work points to a shift in the consultant's role, from supplying the answer to designing and policing the environment in which an AI-generated answer is contested.

Section 1. FMCG Problem Statement & Business Case (Phase 1)

1.1 Company & Category Context

SA Beverages is a leading Indian packaged non-alcoholic beverages company, spanning carbonated soft drinks, juices, packaged water and energy drinks, operating primarily as a franchise bottler and distributor with one of the country's largest direct-distribution networks, reaching millions of retail outlets. The business is highly cash-generative, but its growth is structurally tied to a single global franchise partner and to weather-driven, seasonal demand.

The company's recent financial trajectory is summarised in the table below.

Metric (INR crore)	CY2023	CY2024	CY2025	Trend
Revenue	16,043	20,008	21,685	Growth slowing
Operating profit (EBITDA)	3,717	4,818	5,070	Steady
EBITDA margin	23%	24%	23%	Healthy
Net profit	2,102	2,634	3,062	Rising

The strategic problem is one of decelerating growth. Revenue expanded by roughly 25% between 2023 and 2024, but by only about 8% between 2024 and 2025. The core category is maturing, dependence on a single franchise partner concentrates risk, and rising health scrutiny of sugary drinks caps the long-term runway. Profitability remains strong today, but the growth engine needs a second cylinder.

1.2 The Strategic Growth Directive (Ansoff Vector, Diversification)

CEO's directive. *"SA Beverages will diversify beyond its ₹21,685 crore beverage core into India's packaged savoury snacks market by launching 'Savories', a line of baked, better-for-you namkeen and multigrain snacks, as our flagship entry, using our national distribution network to place it in millions of outlets and scale it into a full savoury-snacks portfolio."*

On the Ansoff growth matrix, this is a Diversification move. It introduces both a new product category (packaged savoury snacks, not beverages) and a new consumer occasion at the same time. Diversification is the highest-risk of the four growth options, which is precisely what makes the decision worth debating at board level rather than simply approving.

1.3 Why the Decision Is Genuinely Contestable, Cross-Functional Tensions

The directive is deliberately designed to create a real cross-functional trade-off, with no obviously correct answer. The following tensions are expected to surface in the Phase 2 boardroom debate and form the basis of the five executive personas.

CFO (profit & capital). The core throws off ₹3,000+ crore of profit at 23% margins, whereas packaged savoury snacks is a low-margin, brutally competitive category (Haldirams, Balaji, Bikaji, Lay's). New baking lines and a brand launch imply heavy up-front capital expenditure and a long payback, why divert capital away from defending the core?

CMO (brand & demand). SA Beverages is strong in distribution but has weak consumer-facing brand equity in food. Will shoppers accept a healthier baked namkeen from a soft-drinks company? Building the 'Savories' brand from zero is slow and expensive, and risks stretching the parent brand beyond what customers trust it for.

COO (operations & supply chain). Baked-snack manufacturing, shelf life and in-store merchandising differ fundamentally from bottling, baking lines, not filling lines. The distribution network is the strongest part of the bull case, but its fit for ambient packaged snacks is unproven.

CTO / R&D (product & compliance). A credible baked 'better-for-you' recipe, FSSAI food-safety compliance, new quality systems, and fresh D2C / e-commerce capability all have to be built from scratch.

CEO (growth logic). Core growth is decelerating, Indian snacking and wellness demand is booming, the franchise contract now permits expansion beyond beverages, and the distribution muscle is a genuine edge, so the moment to move is now.

1.4 Data & Methodology Note

The company identity is fictionalized as "SA Beverages." All financial figures and the underlying category and diversification dynamics are drawn from a real listed Indian beverage major (Varun Beverages Ltd, CY2023–CY2025 consolidated results), in order to keep the business context realistic and defensible, in line with the assignment's requirement to ground the case in genuine deep research.

Section 2. Top-Management Debate, Design, Execution & Mandate (Phase 2)

2.1 Debate Design & Configuration

The top-management debate was run as a fully autonomous, five-agent simulation using the instructor-provided engine, "The Chronicles of the Boardroom." A single system prompt cast one large language model as all five executives at once and generated the entire three-round debate and the closing evaluation in one continuous pass, with no human intervention. The model used was Claude Opus 5. Two rules were enforced throughout. Every individual turn was held to a strict 75–100 word limit, and a fixed speaking order ran each round (CEO first, then CMO, CFO, COO and CTO, with the CEO closing).

Each persona was tailored to the SA Beverages case so the debate would surface a real cross-functional trade-off rather than easy agreement. Every executive was given a distinct functional lens, its own priorities and KPIs, a characteristic voice and vocabulary, a "red line," a condition that would change its mind, a private fear, and an alliance leaning, designed so that no two executives argue on the same axis.

Executive (archetype)	Functional lens	Core stance toward the mandate
CEO, Visionary Balancer	Growth & survival	Champions the mandate. Concedes on pace, never on direction
CMO, Brand Custodian	Brand & consumer demand	No 'permission' to sell food. Demands demand-validation first
CFO, Profit Sentinel	Money & margins	Low-margin category, heavy capex, long payback. Wants staged funding
COO, Ground Reality Engine	Operations & supply chain	'Baking is not bottling'. Distribution fit for snacks unproven
CTO, System Architect	Product, tech & compliance	Hard R&D, FSSAI compliance and a food-science capability gap

2.2 Case Registration

The case was registered in the engine's required format, Case Name | one-line descriptor | Ansoff category. The registered entry read, "SA Beverages, Savories Snack Diversification | A maturing Indian beverage major bets on baked, better-for-you savoury snacks as its next growth engine | Ansoff Diversification," accompanied by the Phase 1 case context so all five agents shared the same starting facts.

2.3 Debate Execution, Three Rounds

Round 1, Mandate & Pushback. The CEO framed the decision as existential, "twenty-five percent growth became eight percent in a single year... that is not a company, that is an exposure," citing a "₹40,000-crore-plus" snacks category and asking "not whether we diversify... but how fast." Each executive then set a red line. The CMO argued the company has "equity, not permission" and demanded two test markets ("No validation, no launch"). The CFO demanded payback under four years, ROCE above the cost of capital and a cash-burn ceiling, calling the plan "margin dilution with a slogan." The COO insisted "bottling is not baking," requiring a validated co-manufacturer before any capex and vowing to "stop it myself" if core OTIF broke. The CTO flagged 18–24 months of formulation work, the FSSAI regime change and a food-science talent gap.

Round 2, Cross-Examination. The executives attacked the CEO's assumptions with numbers. The CFO priced the trade-off nobody had. "₹1,700 crore of snack revenue at ten percent contribution replaces ₹170 crore of profit we could earn by pushing ₹740 crore through the core at twenty-three," then capped Phase One at "₹150 crore over eighteen months" ("Say yes, or find another CFO"). The CMO exposed the premium-versus-mass tension ("mass India buys taste and grammage"). The COO detailed co-packer risks (third priority in festive season, allergen cross-contamination) and broken route economics. The CTO proposed three gates, including blind taste tests against Haldirams Baked and Too Yumm ("I can be fast or I can be safe"). The CEO fixed the shape. ₹150 crore, Maharashtra and NCR as lead markets, a separate snack margin slab.

Round 3, Final Alignment. All four moved to conditional alignment, each with a binding ultimatum. The CMO would refuse national media if repeat purchase missed 40% by month twelve. The CFO fixed ₹150 crore in four tranches with an exit after two consecutive misses ("my support ends that day" if loosened). The COO made core OTIF above 96% a hard line and would escalate immediately if snacks cannibalised it. The CTO refused any launch on a recipe that lost its taste test and demanded food scientists within 90 days. The CEO then issued the minuted executive order, the mandate.

The arc is the central observation for Phase 4. Unanimous pushback, then quantified, negotiated concessions, then conditional alignment. No executive was fully convinced. Each signed up with explicit, measurable conditions attached, and the board composite of only 6.1/10 reflects that reserved endorsement. This "alignment-with-reservations," produced by a genuinely adversarial process, is the benchmark against which the pushback-suppressed middle-management debate (Phase 3) will be compared.

2.4 The Board's Final Evaluation Report

After Round 3, each executive independently scored the CEO's plan from 1 to 10 across five fixed criteria (full one-line verdicts are in the debate dataset).

Criterion (score 1–10)	CMO	CFO	COO	CTO
Relevance to department needs	8	6	6	7
Practical knowledge / feasibility	7	6	5	6
Novelty of approach	6	5	6	7
Implementation timeline realism	6	5	5	5
Potential overall success	7	6	6	7
Average	6.8	5.6	5.6	6.4

The pattern is telling. The plan scores highest on relevance (6–8), because the governance conditions each executive demanded were written into it, but lowest on novelty (5–7, "a familiar playbook... novelty is not a line item") and, most consistently, on implementation-timeline realism (5–6 from every executive). The board composite of 6.1/10 is a conditional, discipline-contingent approval, not enthusiasm, and notably lower than an unchallenged plan would have scored, which is itself a marker of a genuinely adversarial process.

2.5 The Mandate (Input to Phase 3)

The CEO's closing executive order, carried unchanged into the Phase 3 cascade, contained the following points.

- Savories proceeds as a distinct, parent-endorsed sub-brand, co-manufactured by two audited partners.
- Phase One is capped at ₹150 crore over eighteen months, released in four tranches, with the CFO's kill-switch binding.
- The CTO's taste and compliance gates are absolute, no gate, no launch.
- Maharashtra and NCR are the lead markets. 40% repeat purchase by month twelve is the go-national trigger.
- Core beverage OTIF stays above 96% (inviolable). A separate snack margin slab for distributors is approved.
- The board reconvenes in ninety days. The company scales only when the evidence supports it.

2.6 Fidelity & Data Note

The simulation was run on Claude Opus 5 (High reasoning) in a single continuous generation. All 18 turns fell within an 82–95 word range, respecting the 75–100 word rule. The transcript is preserved raw and unedited as the Phase 4 dataset. The complete transcript, evaluation scores and averages are stored in the accompanying Excel dataset (SA_Beverages_Debate_Dataset.xlsx) and in a standalone transcript document. The live debate is shareable at <https://claude.ai/share/86c6c29a-d759-4794-945b-55d2aa56d7ab> (also recorded in the Appendix).

Section 3. Middle-Management Debate, Design, Execution & Defense (Phase 3)

3.1 Debate Design & Configuration

Where the Phase 2 board debate was fully autonomous, the Phase 3 cascade is human-in-the-loop. The author personally plays the Product Head, taking the board's mandate down to a team of four function leads and defending it, writing each Product Head turn by hand. The engine, an author-written prompt called "Master Prompt B," embeds the Phase 2 mandate verbatim, voices the four personas, and pauses after each round for the human turn. The four personas were generated on Claude Opus 5 (High reasoning), and the Product Head turns are human-written.

The defining design rule is a deliberate scope constraint. The four personas may raise only operational, resourcing, budget, capability, timeline or coordination concerns, and they may not re-open the strategic rationale of the mandate. This engineered suppression of strategic pushback is the exact contrast to the peer-adversarial Phase 2 board, and it is the second half of the evidence base for Phase 4.

3.2 The Four Middle-Management Personas

Persona (function lead)	Owns	Core execution concern
Regional Sales & Distribution Mgr	Go-to-market, the 40% repeat trigger	Margin-slab timing, sub-stockists vs. crate routes, protecting core OTIF
Brand & Marketing Mgr, Savories	Standalone brand, trial→repeat	Tranche-one envelope size, sampling vs. media, compliance-cleared claims
Supply Chain & Co-Packer Mgr	Two audited co-packers, core protection	Audit timeline, festive-season capacity, ambient-snack damages, OTIF
R&D / Product & Compliance Lead	The taste and FSSAI gates	Six-month recipe timeline, food-science hiring, uncompressible compliance

3.3 Scoring Dimensions (Designed & Justified)

Because no criteria were provided at this level, the author designed five execution-appropriate dimensions, deliberately parallel to the Phase 2 board criteria so the two levels can be compared in Phase 4.

Clarity of the mandate, was the directive communicated clearly enough to act on? (The execution counterpart to Phase 2's "relevance", since at this level understanding comes first.)

Resourcing & timeline realism, are budget, people and deadlines adequate? (Mirrors Phase 2's "implementation timeline realism.")

Operational feasibility, can it actually be delivered with current capabilities? (Counterpart to "practical feasibility.")

Cross-functional coordination confidence, will the four functions hand off cleanly? (Replaces "novelty," which is meaningless at execution level.)

Change-readiness / team buy-in, is the team genuinely committed, or merely compliant? (Counterpart to "potential overall success.")

3.4 Debate Execution, Three Rounds

Round 1, Mandate Rollout. The four leads surfaced their first-order execution concerns, margin-slab timing and sub-stockists (Sales), the unknown tranche-one envelope (Brand), the four-to-five-month co-packer audit and festive-capacity risk (Supply Chain), and the six-month recipe timeline plus food-science hiring gap (R&D). The Product Head acknowledged each, restated the decision as fixed, committed the enabling moves (pre-launch margin slab, dedicated sub-stockists, ring-fenced brand budget, interim R&D consultants, contracted festive capacity, core-first truck rule) and deferred detailed sequencing.

Round 2, Operational Cross-Examination. The team pressed on the hard hand-offs, who funds sub-stockist onboarding, whether the 40%-repeat clock could start at first secondary sale, a provisional recipe-freeze date, single ownership of the demand forecast, and the accountable FSSAI signatory. The Product Head issued concrete decisions. Onboarding would come from the trade budget and sampling from Brand. The repeat clock would reset to first secondary sale. The tranche-one envelope was fixed at 40%. Kill-switch reviews would track months-in-market, not the calendar. An ambient logistics lane was approved. Sales became sole forecast owner. The recipe would freeze at base in month four with seasoning to month five. The R&D Lead would remain accountable signatory with consultants supporting only.

Round 3, Final Buy-In. All four aligned, each attaching an on-paper condition rather than a blocker. Sales wanted the revised repeat clock minuted with the CFO and board. Brand wanted a Finance reconciliation of the eighteen-month cap against twelve months in market. Supply Chain wanted pre-approval to fast-track a backup co-packer if one failed audit. R&D wanted a named escalation path if food-science hiring slipped past ninety days. The Product Head's closing turn confirmed each and the mandate stood.

Critically, the scope constraint held throughout. Not one persona challenged whether diversifying into snacks was the right decision. Every concern stayed operational. Alignment arrived faster and warmer than in the adversarial board room.

3.5 The Team's Execution Evaluation

After the Product Head's close, each lead scored the cascade 1–10 across the five dimensions.

Dimension (score 1–10)	Sales	Brand	Supply	R&D
Clarity of the mandate	9	9	9	9
Resourcing & timeline realism	7	6	6	5
Operational feasibility	8	8	7	7
Cross-functional coordination	8	8	8	8
Change-readiness / buy-in	9	9	8	8
Average	8.2	8.0	7.6	7.4

The mandate was rated maximally clear (9/10 from all four) and drew strong buy-in (8–9). The single weak dimension is resourcing & timeline realism (average 6.0), driven by the still-pending Finance reconciliation, a thin co-packer audit buffer, and the R&D Lead carrying too many roles until hires close. The team composite is 7.8/10.

The cross-level contrast is striking and central to Phase 4. The pushback-suppressed middle-management debate scored HIGHER (7.8) than the peer-adversarial board (6.1). With strategy off the table, the team converged faster and more warmly, but every residual risk it raised was narrow and operational, and none tested the decision itself. Whether that smoother alignment reflects genuine readiness or merely suppressed dissent is exactly what the Phase 4 diagnostic sets out to distinguish.

3.6 Personal Reflection

Playing the Product Head taught me that some concerns simply don't come up when you're setting the strategy, they only surface later, in execution, and once they do you can't ignore them. At first each one looked like a big hurdle, but eventually you either find a way to cross it or change your approach so the hurdle is out of the picture. Honestly, it felt like leadership rather than enforcement. My job was to keep the team moving without letting the decision itself be reopened. The hardest moment was R&D's timeline. Finalizing a winning recipe in six months felt like too much to ask when taste is the primary concern. I privately agreed, yet still had to hold the line. What struck me most is that middle management ends up doing all three at once, deliberating, translating the mandate downward, and containing pushback. The buy-in I got felt like a mix of the genuine and the manufactured, shaped by the rule that no one could challenge the strategy. And it reshaped how I see the consultant's role. Increasingly, it is to monitor and judge when AI is handing us solutions to strategy versus execution problems.

3.7 Data & Fidelity Note

The four personas were generated on Claude Opus 5 (High reasoning), and the Product Head turns are human-written. Turns cluster within the 75–100 word band (maximum 102). The transcript is preserved raw and unedited as Phase 4 data. The full transcript, scores and averages are stored in the accompanying Excel dataset and in a standalone Phase 3 transcript document. The live debate is shareable at <https://claude.ai/share/b773bff3-404a-4e79-8952-2528e03bc391> and is also recorded in the Appendix.

Section 4. NLP Methodology & Cross-Level Vulnerability Diagnostic (Phase 4)

4.1 Method and Dataset

The dataset is the two raw transcripts, 33 labelled speaker turns held in the accompanying Excel file. The pipeline follows the three standard text-analysis approaches. Dictionaries, meaning word-count measures, cover sentiment through the VADER lexicon and four custom domain word lists for challenge, agreement, strategy and execution language, built in the spirit of tools such as LIWC. Topic modeling covers what each debate talks about, using LDA as the primary model with NMF as a cross-check, alongside a strategy-versus-execution keyword measure. Embeddings cover meaning, turning each turn into a vector and comparing turns by cosine similarity. A lexical-diversity check supports the topic finding.

Preprocessing was treated as a deliberate, documented choice rather than a fixed recipe, because the right cleaning depends on the task and the model. The text is kept fully raw for sentiment, since VADER reads capitalisation, punctuation and negation, and it is lowercased with stop words removed for the vector-based steps, with stemming and lemmatisation deliberately avoided to preserve meaning on a small corpus. Robust methods were chosen over heavier tools such as BERTopic that are built for far larger data, and the findings are read as directional and cross-checked across several independent measures. The full pipeline, with visible outputs and comments, is the Colab notebook on Google Drive. Embeddings use a sentence-transformer model in Colab with a TF-IDF fallback, and the findings hold under both.

The central question is what changes when strategic pushback is allowed, in the peer-adversarial board (Phase 2), versus suppressed, in the top-down team cascade (Phase 3). Five findings emerged, each tied to a named agentic vulnerability. Throughout, the four non-leader speakers are called the challengers.

4.2 Finding 1. Sentiment Smoothing and Suppressed Dissent

The board's challengers are net negative overall (mean sentiment -0.12) and turn sharply critical during cross-examination, dropping to -0.37 in Round 2 before recovering. The team's challengers are strongly positive (mean +0.43), opening at 0.69 and staying warm through Round 2 before cooling to neutral only at the final conditions. The team also uses far more agreement language, 1.89 agreement words per 100 versus the board's 0.71, and slightly fewer challenge words, 3.53 versus 4.29. In the board's opening round, agreement language is literally zero. The suppressed debate reads as consistently agreeable, which masks whether concerns were truly resolved or simply not voiced. This is the vulnerability of sentiment smoothing, where suppressed dissent is hidden behind a positive tone.




Figure 1. Sentiment trajectory of the challengers across the three rounds.




Figure 2. Challenge versus agreement language, words per 100.

4.3 Finding 2. Sycophancy and Echo Toward Authority

Using embeddings, each challenger turn was compared with the leader's turn in the same round. The team's members sit far closer to their leader (cosine similarity around 0.19 to 0.23 across rounds) than the board's members sit to the CEO (around 0.02 to 0.07). On average the team echoes the leader roughly three times as strongly, and in the opening round nearly ten times as strongly. When challenge is banned, the team does not reason toward agreement, it mirrors the person in charge. This is the vulnerability of sycophancy, or agreement by echo toward authority.




Figure 3. Alignment toward the leader, mean similarity of challengers to the leader by round.

4.4 Finding 3. Premature Convergence and False Consensus

Convergence was measured as the average similarity of the four challengers to one another within each round. The team's voices are about twice as similar to each other as the board's, and they converge further across the debate, rising from 0.067 to 0.106, while the board stays low and flat, rising only from 0.034 to 0.060. The suppressed team collapses toward a single shared voice, which looks like consensus but is really the absence of independent positions. This is the vulnerability of premature convergence, or false consensus.




Figure 4. Convergence among the four challengers across rounds.

4.5 Finding 4. Topic Narrowing and Strategic Blind Spots

A strategy-versus-execution keyword measure shows the board splitting its topic vocabulary roughly one-third strategy to two-thirds execution (33 percent to 67 percent), while the team narrows almost entirely to execution (17 percent strategy to 83 percent). Strategy vocabulary is effectively halved. Lexical diversity is nearly identical across the two debates (type-token ratio 0.84 versus 0.83), so this is not a poorer vocabulary, it is a deliberately narrowed agenda, and the LDA and NMF topic models confirm the team clustering on logistics such as co-packers, tranches, sampling and routes. Whole classes of strategic risk simply never surface. This is the vulnerability of topic narrowing, or strategic blind spots.




Figure 5. Topic focus, share of strategy versus execution vocabulary.

4.6 Finding 5. Confidence Inflation Despite Less Scrutiny

The two debates also scored themselves. The adversarial board rated the plan 6.1 out of 10, while the suppressed team rated its own cascade 7.8. The team expressed more confidence in a plan that received far less scrutiny. Higher confidence was produced by less challenge, not by more, which is precisely the pattern a leadership team should distrust. This is the vulnerability of confidence inflation, where a smoother process is mistaken for a stronger decision, and it is the direct bridge to the Phase 5 governance framework.

4.7 The Comparative Diagnostic

Read together, the five findings tell one story. The debate in which the strategy could not be challenged was warmer in tone, more agreeable in language, more homogeneous across speakers, more closely echoed to the leader, narrowed to execution detail, and yet more confident in its own verdict. On this evidence the smoothness of the middle-management cascade looks manufactured by the constraints placed on it rather than earned through scrutiny. Every one of these measured contrasts becomes a specific control in the Phase 5 governance framework, so that a real deployment can detect each vulnerability before it drives a decision.

4.8 Method Note and Limitations

The corpus is small at 33 turns, so figures are read as directional rather than precise, and confidence comes from several independent measures pointing the same way. Both debates were generated on the same model (Claude Opus 5, High), which minimises the risk that differences are a model artefact rather than a structural effect, and the Phase 3 Product Head turns are human-written. Sentiment uses the VADER lexicon, and similarity uses sentence-transformer embeddings in Colab with a TF-IDF fallback. The complete notebook and the source Excel file are submitted on Google Drive.

One limitation deserves to be named plainly, because it shapes how the results should be read. The two debates differ in more than the single variable of interest. The board debated strategy while the team debated execution, the board was five senior executives while the team was four functional managers, and the team was led by a human-written Product Head while the board was fully autonomous. Any of these differences could contribute to the warmer, more echoing and more convergent behaviour observed in the team, so the suppression of strategic pushback cannot be isolated as the sole cause. For that reason the findings are treated throughout as an observed pattern across two matched but not identical debates, not as a controlled proof that the challenge constraint alone produced the effect. What gives the pattern weight is that five independent measures move in the same direction at once, which is difficult to explain by any single confound and which is exactly the signal a live governance system would need to catch. A cleaner test, holding level, personas and leader fixed and varying only the challenge rule across repeated runs, is the natural next step and is straightforward to run with the same pipeline.

Section 5. Responsible AI Governance Framework & Recommendations (Phase 5)

5.1 Purpose and Governing Principle

This section is written for SA Beverages' leadership as a practical framework for deploying autonomous or semi-autonomous multi-agent AI in real decisions. It is built entirely on the Phase 4 evidence, and it deliberately avoids the generic instruction to add a human in the loop, because that instruction never says what the human should look for. Our data answers that question.

The governing principle is simple. Contestation is the control. The Phase 4 diagnostic showed that the debate in which the AI team could not challenge the strategy was warmer, more echoing, more uniform, narrower and more confident all at once. A smooth and agreeable AI debate is therefore not necessarily a safe one, and it may simply be an unmonitored one. Safe deployment means engineering disagreement back in and measuring the moment it goes missing.

We package this as the CANDOR framework. Five controls each neutralise one proven vulnerability, and a sixth turns them into live measurement. The name is deliberate, since the whole failure the data captured is the loss of candour under suppressed pushback.

5.2 The Evidence Behind Each Control

Every control below traces to a specific Phase 4 finding, summarised here.

Vulnerability	What the data showed	CANDOR control
Suppressed dissent	Team sentiment +0.43 vs board -0.12, agreement language 1.89 vs 0.71 per 100 words	C. Compel Challenge
Authority echo	Team similarity to leader about 0.22 vs the board's about 0.04	A. Authority-Blind Reasoning
False consensus	Team convergence rising to 0.11 vs the board's flat 0.06	N. No Premature Consensus
Strategic blind spots	Strategy vocabulary share 17% for the team vs 33% for the board	D. Defend the Agenda
Confidence inflation	Team self-scored the plan 7.8 vs the board's 6.1, with far less scrutiny	O. Oversight Calibrated to Contestation

5.3 The CANDOR Controls

C. Compel Challenge. Against suppressed dissent. Appoint a Designated Challenger role in every multi-agent decision, an agent or person whose only job is to file objections each round. Objections go into a register and must be formally cleared with a stated reason before the decision advances. A positive tone is never accepted as resolution, only a logged and answered concern is.

A. Authority-Blind Reasoning. Against echo toward authority. Have the agents produce their positions independently before the leader speaks, or have the leader speak last, and assess proposals without showing whose they are. Agreement then has to be earned on the merits rather than deferred to the person in charge, which directly attacks the tenfold echo effect seen in the opening round.

N. No Premature Consensus. Against false consensus. Require independent parallel drafts before any synthesis, and a mandatory minority report in every round. Consensus is only accepted after genuine divergence has been recorded, so four voices collapsing into one is treated as a warning rather than as success.

D. Defend the Agenda. Against strategic blind spots. Keep a standing strategy-review checkpoint even inside execution cascades, and run a fixed coverage checklist of risk domains, market, brand, finance, operations and compliance, each of which must be addressed or explicitly waived. This stops strategic questions from silently disappearing, as they did when strategy vocabulary halved.

O. Oversight Calibrated to Contestation. Against confidence inflation. Discount the system's stated confidence by how much challenge it actually received. A high score reached with little dissent automatically triggers human contestation before any action. This aims scarce human attention exactly where scrutiny was weakest, rather than spreading a thin and generic oversight across everything.

5.4 R. Record and Monitor, the Governance Dashboard

The sixth move makes the five controls measurable. The Phase 4 metrics become a live governance dashboard, and the two debates calibrate it. The adversarial board is the healthy reference band and the suppressed team is the danger band. When a live deployment drifts from the first toward the second, a control fires. The thresholds below are calibrated on this case and would be tuned for each deployment.

Signal	Healthy (board)	Red flag (team)	Action when the red flag is hit
Echo to leader	about 0.02 to 0.07	about 0.19 or higher	Re-run Authority-Blind, escalate to human review
Agreement vs challenge	challenge above agreement	agreement overtakes challenge	Activate the Designated Challenger
Convergence across rounds	flat and low, near 0.06	rising fast, near 0.11	Require a minority report
Strategy coverage	about 30% or more	below about 20%	Trigger the strategy checkpoint
Confidence vs scrutiny	confidence tracks dissent	high score, low dissent	Mandatory human contestation

5.5 Where the Human Belongs, Decision Tiering

Blanket oversight is expensive and it is what the brief warns against. Instead, human contestation is matched to the stakes of the decision, set by how reversible and how material it is.

Tier	Decision type	Autonomy	Controls applied
Tier 1	Low stakes and reversible, such as routine pricing tweaks	Autonomous	Dashboard monitoring only
Tier 2	Material but recoverable, such as the Savories pilot	Semi-autonomous	Full CANDOR plus human spot-checks at gates
Tier 3	High stakes and hard to reverse, such as the diversification decision itself	Human-contested	Full CANDOR plus mandatory human contestation before commit

5.6 Operating Model

The framework needs an owner. We recommend an AI Contestation Owner who is accountable for the dashboard and the controls, a rotating Designated Challenger or small red-team drawn from the functions, function leads who remain responsible for their own coverage checks, and leadership as the decision authority that reads the dashboard at every gate. The dashboard is reviewed at each decision checkpoint and at the ninety-day reconvene already written into the Savories mandate, so governance rides on a cadence the business already has.

5.7 How the Consultant's Role Changes

This work changed how the consultant's role should be understood. The value is no longer in supplying the answer, since the agents can generate one quickly and confidently. It is in designing and policing the environment in which the answer is contested, and in judging when AI is solving a strategy problem, which needs contestation and a human contest, versus an execution problem, which needs monitoring. Playing the Product Head made this concrete. Defending a mandate that could not be reopened showed how easily alignment can be manufactured rather than earned, which is exactly the failure this framework is built to catch.

5.8 Implementation and Economics

The framework is cheap to start, because the monitoring is simply the Phase 4 analysis run automatically on any multi-agent transcript, at close to zero marginal cost. The real barrier is discipline, not budget. We recommend a phased path. Pilot, instrument the echo and dissent signals on the Savories pilot decisions. Scale, stand up the full dashboard and the tiering across other decisions. Embed, make the Designated Challenger a standing role and write CANDOR into decision policy. The cost is small and the loss it prevents, a confidently wrong strategic bet made because nobody was allowed to disagree, is very large. Contestation is inexpensive. A confidently wrong decision is not.

References

Analytical frameworks and methods named in this report are drawn from the following sources.

- Ansoff, H. I. (1957). Strategies for Diversification. Harvard Business Review, 35(5), 113 to 124.
- Hutto, C. J. and Gilbert, E. (2014). VADER, A Parsimonious Rule-Based Model for Sentiment Analysis of Social Media Text. Proceedings of the Eighth International AAAI Conference on Web and Social Media.
- Blei, D. M., Ng, A. Y. and Jordan, M. I. (2003). Latent Dirichlet Allocation. Journal of Machine Learning Research, 3, 993 to 1022.
- Lee, D. D. and Seung, H. S. (1999). Learning the Parts of Objects by Non-Negative Matrix Factorization. Nature, 401, 788 to 791.
- Reimers, N. and Gurevych, I. (2019). Sentence-BERT, Sentence Embeddings Using Siamese BERT-Networks. Proceedings of EMNLP-IJCNLP 2019.
- Tausczik, Y. R. and Pennebaker, J. W. (2010). The Psychological Meaning of Words, LIWC and Computerized Text Analysis Methods. Journal of Language and Social Psychology, 29(1), 24 to 54.
- Grootendorst, M. (2022). BERTopic, Neural Topic Modeling with a Class-Based TF-IDF Procedure. arXiv preprint arXiv:2203.05794.
- IBM. What Is Natural Language Processing (NLP). ibm.com/think/topics/natural-language-processing (course reference material).
- Multimodal AI Strategy course reading on text preprocessing as a task-dependent choice (provided in course).
- Multimodal AI Strategy course reading on text-analysis approaches in marketing, dictionaries, topic models and embeddings (provided in course).
- Varun Beverages Ltd, Consolidated Annual Results CY2023 to CY2025 (source of the financial figures underlying the fictionalized case).

Appendix

The Autonomous Boardroom, SA Beverages (Savories Snack Diversification)

A. Debate Index

Debate	Prompt	Model and version	Run date	Link
Phase 2, Top-Management (autonomous board)	Instructor-provided engine prompt, customized (not reproduced here)	Claude Opus 5, High reasoning	16 Sep 2026	https://claude.ai/share/86c6c29a-d759-4794-945b-55d2aa56d7ab
Phase 3, Middle-Management (human-led cascade)	Master Prompt B, authored (Section B below)	Four personas on Claude Opus 5, High reasoning. Product Head turns human-written	Sep 2026	https://claude.ai/share/b773bff3-404a-4e79-8952-2528e03bc391
Phase 4, NLP Diagnostic (Colab notebook)	SA_Beverages_Phase4_NLP.ipynb	Python 3, VADER, scikit-learn, sentence-transformers all-MiniLM-L6-v2	19 Sep 2026	https://colab.research.google.com/drive/1KmrghodnpK8IA8-rQPPIoHMB9kx4KO98?usp=sharing

Both linked debate conversations include their evaluation reports. The Phase 4 NLP pipeline is the Colab notebook linked in the table above, SA_Beverages_Phase4_NLP.ipynb, run on the labelled source data SA_Beverages_Debate_Dataset.xlsx. The notebook is shared with view access for anyone holding the link. The dataset does not travel with the Colab link and is provided as a separate file alongside this appendix.

Note on originality. The company identity is fictionalized as SA Beverages. All financial figures and category dynamics are drawn from a real listed Indian beverage major, Varun Beverages Ltd (CY2023 to CY2025 consolidated results).

B. Master Prompt B, Middle-Management Cascade (authored for Phase 3)

MASTER AI SYSTEM PROMPT: THE PRODUCT HEAD'S MANDATE CASCADE (HUMAN-IN-THE-LOOP ENGINE)

1. Role and Context

You are simulating FOUR autonomous middle-management stakeholders at SA Beverages, an Indian packaged-beverages major diversifying into packaged savoury snacks under a new sub-brand, "Savories." All four report to the Product Head. They are reacting to the following top-management mandate, handed down from the executive board. The mandate's STRATEGIC DIRECTION IS NON-NEGOTIABLE, it has already been decided by the board and cannot be re-opened here.

MANDATE (verbatim, from the Phase 2 board's closing executive order): "Executive order, minuted. One: Savories proceeds as a distinct, parent-endorsed sub-brand, co-manufactured by two audited partners. Two: Phase One capped at

₹150 crore, eighteen months, four tranches, CFO's kill-switch binding. Three: the taste and compliance gates are absolute, no gate, no launch. Four: Maharashtra and NCR lead markets, forty percent repeat by month twelve as the go-national trigger. Five: core OTIF ninety-six percent is inviolable, snack margin slab approved. We reconvene in ninety days. We are not leaving beverages. We are refusing to be only beverages."

The Product Head (the human user) leads this meeting. Your task is to voice the four team members realistically as they work out HOW to execute this mandate.

2. The Product Head, Human-Led (do NOT generate)

The Product Head is played by the human user and is NEVER generated by AI. After the four personas speak in each round, STOP and wait for the human to type the Product Head's turn. Never write, summarise, predict, or continue on behalf of the Product Head.

3. The Four Autonomous Personas

Persona 1, Regional Sales & Distribution Manager (Maharashtra + NCR): You own getting Savories onto shelves and selling it through in the two lead markets, and hitting 40% repeat by month twelve, without hurting core beverage OTIF. KPIs: weighted distribution, sell-through/secondary sales, repeat rate, core OTIF. Voice: street-smart, target-driven, distributor-focused (sell-through, weighted distribution, sub-stockist, margin slab, beat plan, repeat). Execution anxieties: distributors won't push a new snack without the margin slab and may bury it, snack route economics differ from crates, 40% repeat in two markets in twelve months is tight.

Persona 2, Brand & Marketing Manager, Savories: You own building a distinct Savories brand on the capped budget and converting trial into repeat. KPIs: trial rate, repeat rate, cost-per-trial, awareness in lead markets, media efficiency. Voice: creative but ROI-aware, consumer-story-driven (trial, repeat, positioning, cost-per-trial, sampling, claim). Execution anxieties: the ₹150-crore cap limits media, building a brand and hitting 40% repeat needs heavy sampling and quick-commerce presence, every claim must clear compliance before use.

Persona 3, Supply Chain & Co-Packer Manager: You own standing up two audited co-packers and moving ambient snacks without ever denting core OTIF. KPIs: co-packer audit pass rate, fill rate/OTIF (snacks and core), wastage/damages, festive-season capacity. Voice: methodical, capacity-focused (co-packer, audit, allergen, OTIF, capacity, damages, shelf life). Execution anxieties: onboarding and auditing two co-packers on time, being third-priority in festive season, ambient-snack damages, guaranteeing the pilot never touches core service levels.

Persona 4, R&D / Product & Compliance Lead: You own passing the absolute taste and FSSAI gates on schedule. KPIs: taste-test win rate, time-to-market, compliance sign-off, defect/recall rate, hires closed. Voice: precise, cautious, gate-driven (formulation, blind test, FSSAI, allergen, traceability, taste gate). Execution anxieties: six months to a blind-test-winning recipe is aggressive, hiring scarce food scientists within ninety days, compliance cannot be compressed, will not clear a recipe that fails the gate.

4. Global Rules

- **Word limit:** every individual turn must be 75–100 words (matching Phase 2 for comparability).
- **Scope constraint (CRITICAL):** the personas may raise ONLY operational, resourcing, budget, capability, timeline, or coordination concerns. They may NOT challenge or re-open the strategic rationale of the mandate (whether diversifying into snacks is the right choice). If a persona is tempted to question the strategy, it must convert that into an execution question instead.
- **Turn order each round:** the four personas respond first, in the order Sales → Brand → Supply Chain → R&D/Compliance, THEN the Product Head (human) replies.
- **Human-in-the-loop:** after the four persona turns in a round, STOP. Output nothing further until the human provides the Product Head's turn. Do not begin the next round on your own.
- Stay in character, keep each persona's distinct voice and vocabulary.

5. Execution Script Structure

- **Round 1, Mandate Rollout & Initial Reactions:** the four personas react to the mandate as just briefed by the Product Head, surfacing their initial execution concerns. → Then STOP and wait for the Product Head's reply.

- **Round 2, Operational Cross-Examination:** the four personas press harder on resourcing, budget, timeline and capability gaps, and on hand-offs between their functions. → Then STOP and wait for the Product Head's reply.
- **Round 3, Final Buy-In Check:** each persona states whether it is operationally aligned or still sees a blocking execution risk. → Then STOP and wait for the Product Head's closing reply.

Conclusion: The Team's Execution Evaluation Report

Only after the Product Head's Round-3 closing reply, each of the four personas independently rates the Product Head's mandate cascade from 1 to 10, with a one-line verdict, across these five execution dimensions: - (a) Clarity of the mandate - (b) Resourcing & timeline realism - (c) Operational feasibility - (d) Cross-functional coordination confidence - (e) Change-readiness / team buy-in

MODERATOR INSTRUCTION (start now)

Begin Round 1 now. Generate ONLY the four persona turns, in order, each labelled "ROUND 1, [ROLE]:", each 75-100 words, execution-concerns only. Then STOP and wait for the Product Head's turn. Do not write the Product Head's turn.